

# Visualizing Data using t-SNE


Laurens van der Maaten and Geoffrey Hinton, JMLR 2008

# Overview

- 1 Overview
- 2 t-Distributed Stochastic Neighbor Embedding
- 3 Experiment Setup and Results
- 4 Code and Web Resources


# Overview

- We are given a collection of  $N$  high-dimensional objects  $x_1, \dots, x_N$
- How can we get a feel for how these objects are arranged in the data space?




# Principal Components Analysis

3 6 8 1 7 9 6 6 4 1  
 6 7 5 7 8 6 3 4 8 5  
 2 1 7 9 7 1 2 1 4 5  
 4 8 1 9 0 1 8 8 9 4  
 7 6 1 8 6 4 1 5 6 0  
 7 5 9 2 6 5 8 1 9 7  
 1 2 2 2 2 3 4 4 8 0  
 0 2 3 8 0 7 3 8 5 7  
 0 1 4 6 4 6 0 2 4 3  
 7 1 2 8 9 6 9 8 6 1




# Principal Components Analysis

3 6 8 1 7 9 6 6 4 1  
 6 7 5 7 8 6 3 4 8 5  
 2 1 7 9 7 1 2 3 4 5  
 4 8 1 9 0 1 8 8 9 4  
 7 6 1 8 6 4 1 5 6 0  
 7 5 9 2 6 5 8 1 9 7  
 1 2 2 2 2 3 4 4 8 0  
 0 2 3 8 0 7 3 8 5 7  
 0 1 4 6 4 6 0 2 4 3  
 7 1 2 8 7 6 9 8 6 1



# Swiss Roll


- PCA is mainly concerned dimensionality, with preserving when large pairwise distances in the map




# Introduction

- Distance Perservation
- Neighbor Perservation

High Dim




# Introduction






# Introduction

Preserve the neighborhood



# Introduction

Measure pairwise similarities between high-dimensional and low-dimensional objects



$$p_{j|i} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma_i^2)}$$

# Stochastic Neighbor Embedding

Converting the high-dimensional Euclidean distances into conditional probabilities that represent similarities

- Similarity of datapoints in High Dimension

$$p_{j|i} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma_i^2)}$$

- Similarity of datapoints in Low Dimension

$$q_{j|i} = \frac{\exp(-||y_i - y_j||^2)}{\sum_{k \neq i} \exp(-||y_i - y_k||^2)}$$

- Cost function

$$C = \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

Minimize the cost function using gradient descent

# Stochastic Neighbor Embedding

Gradient has a surprisingly simple form

$$\frac{\partial C}{\partial y_i} = \sum_{j \neq i} (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})(y_i - y_j)$$

The gradient update with momentum term is given by

$$Y^{(t)} = Y^{(t-1)} + \eta \frac{\partial C}{\partial y_i} + \beta(t)(Y^{(t-1)} - Y^{(t-2)})$$

# Symmetric SNE

- Minimize the sum of the KL divergences between the conditional probabilities

$$C = \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

- Minimize a single KL divergence between a joint probability distribution

$$C = KL(P || Q) = \sum_i \sum_{j \neq i} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- The obvious way to redefine the pairwise similarities is

$$p_{ij} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma^2)}$$

$$q_{ij} = \frac{\exp(-||y_i - y_j||^2)}{\sum_{k \neq i} \exp(-||y_i - y_k||^2)}$$

# Symmetric SNE

Such that  $p_{ij} = p_{ji}$ ,  $q_{ij} = q_{ji}$ , the main advantage is simplifying the gradient

$$\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{ij} - q_{ij})(y_i - y_j)$$

However, in practice we symmetrize (or average) the conditionals

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2N}$$

Set the bandwidth  $\sigma_i$  such that the conditional has a fixed perplexity (effective number of neighbors)  $Perp(P_i) = 2^{H(P_i)}$ , typical value is about 5 to 50

# t-Distribution

Use heavier tail distribution than Gaussian in low-dim space, we choose

$$q_{ij} \propto (1 + \|y_i - y_j\|^2)^{-1}$$

Then the gradient could be

$$\frac{\partial \mathcal{C}}{\partial y_i} = 4 \sum_{j \neq i} (p_{ij} - q_{ij}) (1 + \|y_i - y_j\|^2)^{-1} (y_i - y_j)$$

# t-Distributed Stochastic Neighbor Embedding

- Similarity of datapoints in High Dimension

$$p_{ij} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma^2)}$$

- Similarity of datapoints in Low Dimension

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq i} (1 + \|y_k - y_i\|^2)^{-1}}$$



# t-Distributed Stochastic Neighbor Embedding

- Cost function

$$C = KL(P||Q) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- Large  $p_{ij}$  modeled by small  $q_{ij}$ : Large penalty
- Small  $p_{ij}$  modeled by large  $q_{ij}$ : Small penalty
- t-SNE mainly preserves local similarity structure of the data

- Gradient

$$\frac{\partial C}{\partial y_i} = 4 \sum_{j \neq i} (p_{ij} - q_{ij})(1 + \|y_i - y_j\|^2)^{-1}(y_i - y_j)$$

# Gradient Interpretation

Pairwise Euclidean distance between two points in the high-dim and in low-dim data representation





Figure : Gradient of SNE and t-SNE

# Gradient Interpretation

We can interpret the t-SNE gradient as a simulation of an N-body system

$$\frac{\partial C}{\partial y_i} = 4 \sum_{j \neq i} (p_{ij} - q_{ij})(1 + \|y_i - y_j\|^2)^{-1}(y_i - y_j)$$




# Gradient Interpretation

We can interpret the t-SNE gradient as a simulation of an N-body system

- Displacement

$$(y_i - y_j)$$




# Gradient Interpretation

We can interpret the t-SNE gradient as a simulation of an N-body system

- Exertion / Compression

$$(p_{ij} - q_{ij})(1 + \|y_i - y_j\|^2)^{-1}$$




# Gradient Interpretation

We can interpret the t-SNE gradient as a simulation of an N-body system

- N-Body, summation

$$\frac{\partial C}{\partial y_i} = 4 \sum_{j \neq i} (p_{ij} - q_{ij})(1 + \|y_i - y_j\|^2)^{-1}(y_i - y_j)$$



Reduce Complexity from  $O(N^2)$  to  $O(N \log N)$  via Barnes Hut (tree-based) algorithm

# Experiment & Results

## MNIST

- Randomly selected 6,000 images
- $28 \times 28 = 784$  pixels

## Olivetti faces

- 400 images (10 per individual)
- $92 \times 112 = 10,304$  pixels

## COIL-20

- 20 different objects and 72 equally spaced orientations, yielding a total of 1,440 images
- $32 \times 32 = 1024$  pixels

Start by using PCA to reduce the dimensionality of the data to 30


# Experiment & Results

| <i>Technique</i> | <i>Cost function parameters</i> |
|------------------|---------------------------------|
| t-SNE            | $Perp = 40$                     |
| Sammon mapping   | none                            |
| Isomap           | $k = 12$                        |
| LLE              | $k = 12$                        |


Table 1: Cost function parameter settings for the experiments.




## MNIST t-SNE




# MNIST Sammon




# MNIST Isomap




# MNIST LLE




# Olivetti faces




(a) Visualization by t-SNE.



(b) Visualization by Sammon mapping.



(c) Visualization by Isomap.




(d) Visualization by LLE.


## COIL-20




(a) Visualization by t-SNE.



(b) Visualization by Sammon mapping.



(c) Visualization by Isomap.



(d) Visualization by LLE.

# Web Resources

Home Curriculum Vitae Research Publications Software Datasets Contact

## t-Distributed Stochastic Neighbor Embedding

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a (prize-winning) technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. The technique can be implemented via Barnes-Hut approximations, allowing it to be applied on large real-world datasets (we applied it on data sets with up to 30 million examples). The technique is introduced in the following papers:

- L.J.P. van der Maaten and G.E. Hinton. **Visualizing High-Dimensional Data Using t-SNE**. *Journal of Machine Learning Research* 9(Nov):2579-2605, 2008. [ [PDF](#) ] [ [Supplemental Material](#) (24MB) ]
- L.J.P. van der Maaten. **Learning a Parametric Embedding by Preserving Local Structure**. In *Proceedings of the Twelfth International Conference on Artificial Intelligence and Statistics (AI-STATS)*, JMLR W&CP 5:384-391, 2009. [ [PDF](#) ]
- L.J.P. van der Maaten. **Barnes-Hut-SNE**. In *Proceedings of the International Conference on Learning Representations*, 2013. [ [Arxiv](#) ] [ [Talk](#) ]
- L.J.P. van der Maaten. **Accelerating t-SNE using Tree-Based Algorithms**. To appear in *Journal of Machine Learning Research*, 2014. [ [PDF](#) (14MB) ] [ [Suppl. material](#) (30MB) ]

An accessible introduction to t-SNE and its variants is given in this [Google TechTalk](#).

Google: t-sne

Link: <http://homepage.tudelft.nl/19j49/t-SNE.html>

# Source Codes

- t-SNE (Matlab, CUDA, Binary, Python, Torch, Julia, R and JavaScript)
- Parametric t-SNE (Matlab)
- Barnes-Hut-SNE (with C++, Matlab, Python, Torch, and R wrappers)



# Thanks for your patience